

Laboratory Exercise 2

An Enhanced Processor

In Laboratory Exercise 1 we described a simple processor. In Part I of that exercise the processor itself was designed, and in Part II the processor was connected to an external counter and a memory unit. This exercise describes subsequent parts of the processor design. Note that the numbering of figures and tables in this exercise are continued from those in Parts I and II of the preceding lab exercise.

In this exercise we will extend the capability of the processor so that the external counter is no longer needed, and so that the processor has the ability to perform read and write operations using memory or other devices. We will add four new instructions to the processor, as displayed in Table 3. The *ld* (load) instruction loads data into register *rX* from the external memory address specified in register *rY*. The *st* (store) instruction stores the data contained in register *rX* into the memory address found in *rY*. Finally, the instructions *mvnz* (move if not zero) performs a *mv* operation only under the condition that the current contents of register *G* (the adder/subtractor output) is not equal to 0, and the *mvnc* instruction performs a *mv* operation only if the most-recently executed *add* (or *sub*) instruction did not produce a carry-out (or borrow).

Operation	Function performed
<i>ld rX, [rY]</i>	$rX \leftarrow [rY]$
<i>st rX, [rY]</i>	$[rY] \leftarrow rX$
<i>mvnz rX, rY</i>	if $G \neq 0$, $rX \leftarrow rY$
<i>mvnc rX, rY</i>	if no carry-out, $rX \leftarrow rY$

Table 3: New instructions performed in the processor.

A schematic of the enhanced processor is given in Figure 12. All general-purpose registers are now 16-bits wide, whereas in Laboratory Exercise 1 they were nine-bits wide. In the figure registers *r0* to *r6* are the same as in Figure 1 of Exercise 1, but register *r7* has been changed to a counter. This counter is used to provide the addresses in the memory from which the processor's instructions are read; in the preceding lab exercise, a counter external to the processor was used for this purpose. We will refer to *r7* as the processor's *program counter (pc)*, because this terminology is common for real processors available in the industry. When the processor is reset, *pc* is set to address 0. At the start of each instruction (in time step T_0) the contents of *pc* is used as an address to read an instruction from the memory. The instruction is stored into *IR* and the *pc* is automatically incremented to point to the next instruction (in the case of *mvi* the *pc* provides the address of the immediate data and is then incremented again).

The processor's control unit increments *pc* by using the *incr_pc* signal, which is just an enable on this counter. It is also possible to directly load an address into *pc* (*r7*) by having the processor execute a *mv* or *mvi* instruction in which the destination register is specified as *r7*. In this case the control unit uses the signal $r7_{in}$ to perform a parallel load of the counter. Thus, the processor can execute instructions at any address in the memory, as opposed to only being able to execute instructions that are stored in successive addresses. Similarly, the current contents of *pc*, which always has the address of the *next* instruction to be executed, can be copied into another register by using a *mv* instruction.

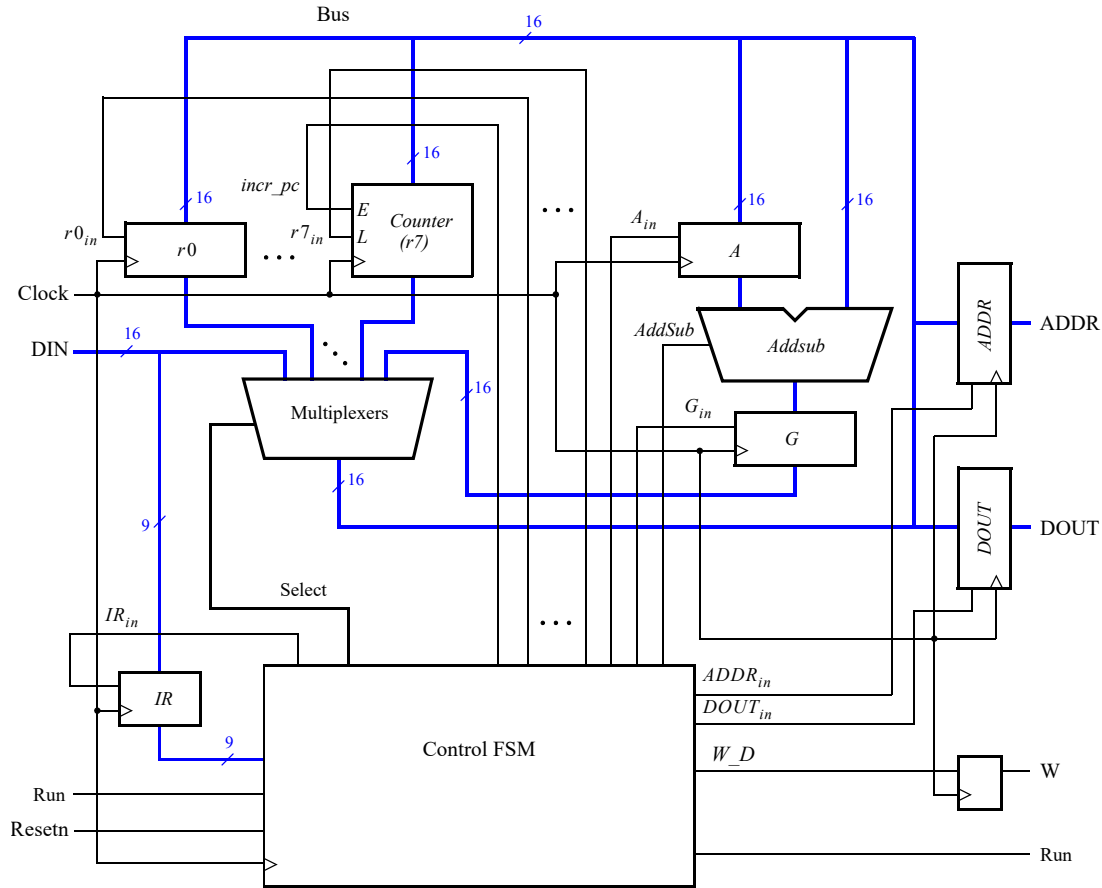


Figure 12: An enhanced version of the processor.

An example of code that uses the *pc* register to implement a loop is shown below, where the text after the *//* on each line is just a comment. The instruction *mv r5,r7* places into *r5* the address in memory of the instruction *sub r2,r1*. Then, the instruction *mvnz r7,r5* causes the *sub* instruction to be executed repeatedly until *r2* becomes 0. This type of loop could be used in a larger program as a way of creating a delay.

```

mvi    r1, #1
mvi    r2, #10000 // delay value
mv     r5, r7     // save address of next instruction
sub    r2, r1     // decrement delay count
mvnz   r7, r5     // loop until delay count gets to 0

```

Figure 12 shows two registers in the processor that are used for data transfers. The *ADDR* register is used to send addresses to an external device, such as a memory module, and the *DOUT* register is used by the processor to provide data that can be stored outside the processor. One use of the *ADDR* register is for reading, or *fetching*, instructions from memory; when the processor wants to fetch an instruction, the contents of *pc* (*r7*) are transferred across the bus and loaded into *ADDR*. This address is provided to the memory. In addition to fetching instructions, the processor can read data at any address by using the *ADDR* register. Both data and instructions are read into the processor on the *DIN* input port. The processor can write data for storage at an external address by placing this address into the *ADDR* register, placing the data to be stored into its *DOUT* register, and asserting the output of the *W* (write) flip-flop to 1.

Figure 13 illustrates how the enhanced processor can be connected to memory and other devices. The memory unit in the figure supports both read and write operations and therefore has both address and data inputs, as well as a write enable input. The memory also has a clock input, because the address, data, and write enable inputs must be loaded into the memory on an active clock edge. This type of memory unit is usually called a *synchronous static random access memory (synchronous SRAM)*. Figure 13 also includes a 9-bit register that can be used to store data from the processor; this register might be connected to a set of LEDs to allow display of data on your DE-series board. To allow the processor to select either the memory unit or register when performing a write operation, the circuit shows NOR gates that perform *address decoding*: if the upper address lines are $A_{15} \dots A_{12} = 0$, then the memory module will be written. Figure 13 shows n lower address lines connected to the memory; for this exercise a memory with 256 words is probably sufficient, which implies that $n = 8$ and the memory address port is driven by $A_7 \dots A_0$. For addresses in which $A_{15} \dots A_{12} = 1$, the data written by the processor is loaded into the register whose outputs are called *LEDs* in Figure 13.

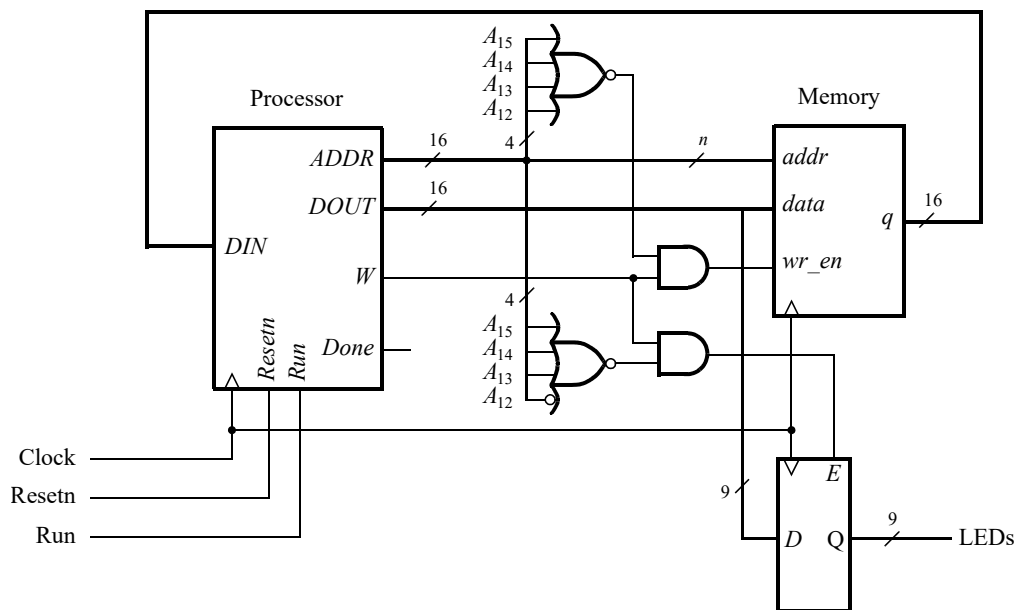


Figure 13: Connecting the enhanced processor to a memory unit and output register.

Part III

Figure 14 gives Verilog code for a top-level file that you can use for this part of the exercise. This file implements the circuit illustrated in Figure 13, plus an additional input port that is connected to switches $SW_8 \dots SW_0$. This input port can be read by the processor at addresses in which $A_{15} \dots A_{12} = 3$. Switch SW_9 is not a part of the switch port, because it is dedicated for use as the processor's *Run* input.

The code in Figure 14 is provided on the course website, along with a few other source-code files: *flipflop.v*, *inst_mem.v*, *inst_mem.mif*, and *proc.v*. The *inst_mem.v* source-code file was created by using the Quartus IP Catalog to instantiate a RAM: 1-PORT memory module. The memory has one 16-bit wide read/write data port and is 256 words deep.

The *proc.v* file provides partially-completed Verilog code for the enhanced processor. This code implements register *r7* as a program counter, as discussed above, and includes a number of changes that are needed to support the new *ld*, *st*, *mvnz*, and *mvnc* instructions. In this part you are to augment the provided Verilog code to complete the implementation of the *ld* and *st* instructions. You do not need to work on the *mvnz* or *mvnc* instructions for this part.

```

module part3 (KEY, SW, CLOCK_50, LEDR);
    input [0:0] KEY;
    input [9:0] SW;
    input CLOCK_50;
    output [9:0] LEDR;

    wire [15:0] DOUT, ADDR;
    wire Done;
    reg [15:0] DIN;
    wire W, Sync, Run;
    wire inst_mem_cs, SW_cs, LED_reg_cs;
    wire [15:0] inst_mem_q;
    wire [8:0] LED_reg, SW_reg;    // LED[9] and SW[9] are used for Run

    // synchronize the Run input
    flipflop U1 (SW[9], KEY[0], CLOCK_50, Sync);
    flipflop U2 (Sync, KEY[0], CLOCK_50, Run);

    proc U3 (DIN, KEY[0], CLOCK_50, Run, DOUT, ADDR, W, Done);

    assign inst_mem_cs = (ADDR[15:12] == 4'h0);
    assign LED_reg_cs = (ADDR[15:12] == 4'h1);
    assign SW_cs = (ADDR[15:12] == 4'h3);
    inst_mem U4 (ADDR[7:0], CLOCK_50, DOUT, inst_mem_cs & W, inst_mem_q);

    always @ (*)
    if (inst_mem_cs == 1'b1)
        DIN = inst_mem_q;
    else if (SW_cs == 1'b1)
        DIN = {7'b0000000, SW_reg};
    else
        DIN = 16'bxxxxxxxxxxxxxxxx;

    regn #(n(9)) U5 (DOUT[8:0], LED_reg_cs & W, CLOCK_50, LED_reg);
    assign LEDR[8:0] = LED_reg;
    assign LEDR[9] = Run;

    regn #(n(9)) U7 (SW[8:0], 1'b1, CLOCK_50, SW_reg); // SW[9] is used for Run
endmodule

```

Figure 14: Verilog code for the top-level file.

Perform the following:

1. Augment the code in *proc.v* so that the enhanced processor fully implements the *ld* and *st* instructions. Test your Verilog code by using the ModelSim simulator. Include in the project the top-level file from Figure 14 and the other related source-code files. Sample setup files for ModelSim are provided on the course website, along with the other files for this exercise. The memory module *inst_mem* will be initialized with the contents of the file *inst_mem.mif*. This file represents the assembly-language program shown in Figure 15, which tests the *ld* and *st* instructions by reading the values of the SW switches and writing these values to the LEDs, in an endless loop. Make sure that your testbench sets the *Run* switch, *SW₉*, so that your processor will execute instructions. Observe the signals inside your processor, toggle the other SW switches, and observe the LED signals in the top-level module.

The file *inst_mem.mif* was created by *assembling* the program in Figure 15 using the Assembler *sbasm.py*. This Assembler is written in Python and is available on the course website. To use this Assembler you have to first install Python (version 2.7) on your computer. The Assembler includes a README file that explains

its use. The *sbasm.py* assembler allows you to define symbolic constants using a `.define` directive, and it supports the use of labels for referring to lines of code, such as `MAIN` in the figure.

```
.define LEDR_ADDRESS 0x1000
.define SW_ADDRESS 0x3000

// Read SW switches and display on LEDs
    mvi    r3, #LEDR_ADDRESS // point to LED port
    mvi    r4, #SW_ADDRESS   // point to SW port
MAIN:   ld     r0, [r4]        // read SW values
        st     r0, [r3]        // light up LEDs
        mvi    r7, #MAIN
```

Figure 15: Assembly-language program that uses *ld* and *st* instructions.

2. Once your simulation results are correct, create a new Quartus project for your Verilog code. Compile your code using the Quartus Prime software, and download the resulting circuit into the DE1-SoC board. Toggle the SW switches and observe the LEDs to test your circuit.

Part IV

In this part you are to create a new Verilog module to be instantiated in your top-level module from Part III. The new module will function as an output port called *seg7_scroll*. It will allow your processor to write data to each of the six 7-segment displays on a DE-series board. Include six write-only seven-bit registers in the *seg7_scroll* module, one for each display. Each register should directly drive the segment lights for one seven-segment display, so that the processor can write characters onto the displays.

Create the necessary address decoding to allow the processor to write to the registers in the *seg7_scroll* module at addresses in which $A_{15} \dots A_{12} = 2$. Address 0×2000 should select the register that controls display *HEX0*, 0×2001 should select the register for *HEX1*, and so on. If your processor writes 0 to address 0×2000 , then the *seg7_scroll* module should turn off all of the segment-lights in the *HEX0* display; writing $0 \times 7f$ should turn on all of the lights in this display. You will need to add ports to your top-level design for the six 7-segment displays. Pin assignments for these ports, which are called *HEX0*[6:0], *HEX1*[6:0], ..., *HEX6*[6:0], are included in the Quartus settings files provided on the course website. Segment 0 is on the top of a seven-segment display, and then segments 1 to 5 are assigned in a clockwise fashion, with segment 6 being in the middle of the display. Thus, the Verilog assignment statement

```
HEX0[6:0] = 7'b01101101;
```

would display the digit **5** on *HEX0*. Perform the following:

1. You may want to first simulate/debug your code with ModelSim. Then, create a Quartus project for this part of the exercise and write the Verilog code for the *seg7_scroll* module. Modify your top-level module to connect the processor to the *seg7_scroll* module, and to connect this module to the HEX display pins.
2. To test your circuit you can use the *inst_mem.mif* file provided for this part of the exercise on the course website, which corresponds to the assembly-language program shown in Figure 16. This program works as follows: it reads the SW switch port and lights up a seven-segment display corresponding to the value read. For example, if the SW switches were set to 0, then the digit **0** would be shown on *HEX0*. If the switches were set to 1, then the digit **1** would be displayed on *HEX1*, and so on, up to the digit **5** which would be shown on *HEX5*.

```

.define LED_ADDRESS 0x1000
.define HEX_ADDRESS 0x2000
.define SW_ADDRESS 0x3000

// This program shows a decimal digit on the HEX displays
    mv     r5, r7           // return address for subroutine
    mvi    r7, #BLANK      // call subroutine to clear HEX displays
MAIN:    mvi    r2, #HEX_ADDRESS // point to HEX port
        mvi    r3, #DATA    // used to get 7-segment display patterns

        mvi    r4, #SW_ADDRESS // point to SW port
        ld     r0, [r4]      // read switches
        mvi    r4, #LED_ADDRESS // point to LED port
        st     r0, [r4]      // light up LEDs
        add    r2, r0        // point to correct HEX display
        add    r3, r0        // point to correct 7-segment pattern

        ld     r0, [r3]      // load the 7-segment pattern
        st     r0, [r2]      // light up HEX display

        mvi    r7, #MAIN

// Subroutine BLANK
// This subroutine clears all of the HEX displays
// input: none
// returns: nothing
BLANK:
    mvi    r0, #0           // used for clearing
    mvi    r1, #1           // used to add/sub 1
    mvi    r2, #HEX_ADDRESS // point to HEX displays
    st     r0, [r2]         // clear HEX0
    add    r2, r1
    st     r0, [r2]         // clear HEX1
    add    r2, r1
    st     r0, [r2]         // clear HEX2
    add    r2, r1
    st     r0, [r2]         // clear HEX3
    add    r2, r1
    st     r0, [r2]         // clear HEX4
    add    r2, r1
    st     r0, [r2]         // clear HEX5

    add    r5, r1
    add    r5, r1
    mv     r7, r5           // return from subroutine

DATA:    .word 0b00111111    // '0'
        .word 0b00000110    // '1'
        .word 0b01011011    // '2'
        .word 0b01001111    // '3'
        .word 0b01100110    // '4'
        .word 0b01101101    // '5'

```

Figure 16: Assembly-language program that tests the seven-segment displays.

Part V

In this part you are to enhance your processor so that it implements the *mvnz* and *mvnc* instructions. To do this, you should create two *condition-code flags* in your processor. One flag, *z*, should be set to 1 when the adder/subtractor unit in the processor generates a result of zero; otherwise *z* should be 0. The other flag, *c*, should be set to 1 when the adder/subtractor unit produces a carry-out of 1; otherwise *c* should be 0. Thus, *c* should be 1 when an *add* instruction generates a carry-out, or when a *sub* instruction requires a borrow for the most-significant bit. The *mvnz* and *mvnc* instructions should perform a *mv* operation when *z* = 0 and *c* = 0, respectively.

The *mv* instruction should also affect the *z* flag, based on whether the result produced by the instruction is zero. For example, if register *r0* contains 0, then *mv r1, r0* should set the *z* flag to 1, because the instruction produces the result *r1* = 0. The *mv* instruction may have an undefined effect on the *c* flag.

Perform the following:

1. Enhance your processor so that it implements the condition-code flags *z* and *c*, and supports the *mvnz* and *mvnc* instructions. Also, modify your processor's *mv* instruction so that it sets, or clears, the *z* flag appropriately. You may want to use ModelSim to test/debug your code, and then create a new Quartus project to implement this part of the exercise in the DE1-SoC board.
2. To test your processor, you can use the *inst_mem.mif* file that is provided for this part of the exercise on the course website. This file corresponds to the assembly-language program displayed in Figure 17. It provides code that tests for the correct operation of instructions supported by the enhanced processor. If all of the tests pass, then the program shows the word **PASSEd** on the seven-segment displays. It also shows the binary value 1000 on the LEDs, representing eight successful tests. If any test fails, then the program shows the word **FAILEd** on the seven-segment displays and places on the LEDs the address in the memory of the instruction that caused the failure. If a failure occurs, then the offending instruction can be identified by cross-referencing the LED pattern with the addresses in the file *inst_mem.mif*.

```
.define LED_ADDRESS 0x1000
.define HEX_ADDRESS 0x2000

// This code tests various instructions
    mvi    r0, #0
    mvi    r1, #1
    mvi    r2, #0        // used to count number of successful tests
    mvi    r5, #FAIL

    mvi    r6, #T1       // save address of next instruction
    add    r0, r0        // set the z flag
// test mvnz
T1:    mvnz r7, r5        // should not take the branch!

    add    r2, r1        // incr success count
    mvi    r4, #S1
    mvi    r6, #T2       // save address of next instruction
// test mv's effect on z flag
    mv     r3, r1        // reset the z flag
    mvnz   r7, r4        // should take the branch!
T2:    mvi    r7, #FAIL

S1:    add    r2, r1        // incr success count
    mvi    r4, #S2
    mvi    r6, #T3       // save address of next instruction
```

Figure 17: Assembly-language program that tests various instructions. (Part a)

```

        add    r3, r1        // reset the z flag
        mvnz   r7, r4        // should take the branch!
T3:     mvi     r7, #FAIL

S2:     add    r2, r1        // incr success count
        mvi     r6, #T4
        mv      r0, r0        // set the z flag
// test mv's effect on z flag
T4:     mvnz   r7, r5        // should not take the branch!

        add    r2, r1        // incr success count
        mvi     r6, #T5        // save address of next instruction
        mvi     r3, #0xffff
        add    r3, r1        // set the c flag
// test mvnc
T5:     mvnc   r7, r5        // should not take the branch!

        add    r2, r1        // incr success count
        mvi     r4, #S3
        mvi     r6, #T6
        add    r3, r0        // clear carry flag
// test mvnc
        mvnc   r7, r4        // should take the branch!
T6:     mvi     r7, #FAIL

S3:     add    r2, r1        // count the success

// finally, test ld and st from/to memory
        mvi     r6, #T7        // save address of next instruction
        mvi     r4, #_LDTEST
        ld      r4, [r4]
        mvi     r3, #0xA5A5
        sub     r3, r4
T7:     mvnz   r7, r5        // should not take the branch!

        add    r2, r1        // incr success count
        mvi     r6, #T8        // save address of next instruction
        mvi     r3, #0xA5A5
        mvi     r4, #_STTEST
        st      r3, [r4]
        ld      r4, [r4]
        sub     r3, r4
T8:     mvnz   r7, r5        // should not take the branch!

        add    r2, r1        // incr success count
        mvi     r7, #PASS

// Show the result on the HEX displays
FAIL:   mvi     r3, #LED_ADDRESS
        st      r6, [r3]        // show address of failed test on LEDs
        mvi     r5, #_FAIL
        mvi     r7, #PRINT
PASS:   mvi     r3, #LED_ADDRESS
        st      r2, [r3]        // show success count on LEDs
        mvi     r5, #_PASS

```

Figure 17: Assembly-language program that tests various instructions. (Part b)


```

PRINT:  mvi    r1, #1
        mvi    r4, #HEX_ADDRESS // address of HEX0
        // We would normally use a loop counting down from 6
        // with mvnz to display the six letters. But in this
        // testing code we can't assume that mvnz even works!
        ld     r3, [r5] // get letter
        st     r3, [r4] // send to HEX display
        add    r5, r1    // ++increment character pointer
        add    r4, r1    // point to next HEX display
        ld     r3, [r5] // get letter
        st     r3, [r4] // send to HEX display
        add    r5, r1    // ++increment character pointer
        add    r4, r1    // point to next HEX display
        ld     r3, [r5] // get letter
        st     r3, [r4] // send to HEX display
        add    r5, r1    // ++increment character pointer
        add    r4, r1    // point to next HEX display
        ld     r3, [r5] // get letter
        st     r3, [r4] // send to HEX display
        add    r5, r1    // ++increment character pointer
        add    r4, r1    // point to next HEX display
        ld     r3, [r5] // get letter
        st     r3, [r4] // send to HEX display
        add    r5, r1    // ++increment character pointer
        add    r4, r1    // point to next HEX display

HERE:   mvi    r7, #HERE

_PASS:  .word  0b0000000001011110 // d
        .word  0b0000000001111001 // E
        .word  0b0000000001101101 // S
        .word  0b0000000001101101 // S
        .word  0b0000000001110111 // A
        .word  0b0000000001110011 // P
_FAIL:  .word  0b0000000001011110 // d
        .word  0b0000000001111001 // E
        .word  0b0000000001110000 // L
        .word  0b0000000001100000 // I
        .word  0b0000000001110111 // A
        .word  0b0000000001110001 // F
_LDTEST: .word  0xA5A5
_STTEST: .word  0x5A5A

```

Figure 17: Assembly-language program that tests various instructions. (Part c)

Part VI

Write an assembly-language program that displays a binary counter on the LED port. Initialize the counter to 0, and then increment the counter by one in an endless loop. You should be able to control the speed at which the counter is incremented by using nested delay loops, along with the SW switches. If the SW switches are set to their maximum value, 0b11111111, then the delay loops should cause the counter to increment slowly enough so that each change in the counter can be visually observed on the LEDs. Lowering the value of the SW switches should make the counter increment more quickly up to some maximum.

Assemble your program by using the *sbasml.py* Assembler (or assemble it by hand if you prefer!). Save the output produced by *sbasml.py* in a new file named *inst_mem.mif*. Make sure that the new *MIF* file is stored in the folder that holds your Quartus project for Part V. Then use the Quartus command Processing > Update Memory Initialization File, to include your new *inst_mem.mif* file in your Quartus project. Next, select the Quartus command Processing > Start > Start Assembler to produce a new programming *bitstream* for your DE-series board. Finally, use the Quartus Programmer to download the new bitstream onto your board. If the *Run* signal is asserted, then your processor will execute your the program.

Part VII

Augment your assembly-language program from Part VI so that counter values are displayed on the seven-segment display port rather than on the LED port. You should display the counter values as decimal numbers from **0** to **65535**. The speed of counting should be controllable using the SW switches in the same way as for Part VI. As part of your solution you may want to make use of the code shown in Figure 18. This code provides a subroutine that divides the number in register r0 by 10, returning the quotient in r2 and the remainder in r0. Dividing by 10 is a useful operation when performing binary-to-decimal conversion.

```
// subroutine DIV10
// This subroutine divides the number in r0 by 10
// The algorithm subtracts 10 from r0 until r0 < 10, and keeps count in r2
// input: r0
// returns: quotient Q in r2, remainder R in r0
DIV10:
    mvi    r1, #1
    sub    r6, r1           // save registers that are modified
    st     r3, [r6]
    sub    r6, r1
    st     r4, [r6]         // end of register saving

    mvi    r2, #0           // init Q
    mvi    r3, RETDIV       // for branching
DLOOP:   mvi    r4, #9       // check if r0 is < 10 yet
    sub    r4, r0
    mvnc   r7, r3           // if so, then return

    INC:   add    r2, r1     // but if not, then increment Q
    mvi    r4, #10
    sub    r0, r4           // r0 -= 10
    mvi    r7, DLOOP        // continue loop

    RETDIV: ld     r4, [r6]   // restore saved regs
    add    r6, r1
    ld     r3, [r6]         // restore the return address
    add    r6, r1
    add    r5, r1           // adjust the return address by 2
    add    r5, r1
    mv     r7, r5           // return results
```

Figure 18: A subroutine that divides by 10

As described for Part VI, assemble your program with *sbasml.py*, update your *MIF* file in the Quartus software, generate a new bitstream file by using the Quartus Assembler, and then download the new bitstream onto your DE-series board to run your code.